

## Part 1

```
In [1]: '''  
1. Identify each column as nominal, ordinal, interval, or ratio in the Auto_mpg_raw.csv data set.  
a.      Miles per gallon: ratio  
b.      Cylinders: ratio  
c.      Displacement: nominal  
d.      Horsepower: ratio  
e.      Weight: ratio  
f.      Acceleration: nominal  
g.      Model Year: interval  
h.      Origin: nominal  
i.      Car Name: nominal  
'''
```

```
Out[1]: '\n1. Identify each column as nominal, ordinal, interval, or ratio in the Auto_mpg_raw.csv data set.\na.\tMiles per gallon: ratio\nb.\tCylinders: ratio\nc.\tDisplacement: nominal\nd.\tHorsepower: ratio\ne.\tWeight: ratio\nf.\tAcceleration: nominal\ng.\tModel Year: interval\nh.\tOrigin: nominal\ni.\tCar Name: nominal\n\n'
```

## Part 2

```
In [2]: # Importing the necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
pd.set_option('display.max_rows', 500)

# importing the Auto_mpg_raw.csv file using pandas
auto_Mpg = pd.read_csv("Auto_mpg_raw.csv")

# Plotting all the histograms
# histogram for miles per gallon
plt.figure()
plt.hist(auto_Mpg["Miles per gallon"])
plt.title("Miles per gallon")

#histogram for Cylinders
plt.figure()
plt.hist(auto_Mpg["Cylinders"])
plt.title("Cylinders")

#histogram for Displacement
plt.figure()
plt.hist(auto_Mpg["Displacement"])
plt.title("Displacement")

# histogram for Horsepower
plt.figure()
plt.hist(auto_Mpg["Horsepower"])
plt.title("Horsepower")

#histogram for Weight
plt.figure()
plt.hist(auto_Mpg["Weight"])
plt.title("Weight")

#histogram for Acceleration
plt.figure()
plt.hist(auto_Mpg["Acceleration"])
plt.title("Acceleration")

#histogram for Model Year
plt.figure()
plt.hist(auto_Mpg["Model year"])
plt.title("Model year")

#histogram for origin
plt.figure()
plt.hist(auto_Mpg["Origin"])
plt.title("Origin")
```

""" Answer to question number 2

Here from seeing the distribution in the histograms, some misleading datapoints are observed.

The misleading data points are discussed below:

1. The first attribute Miles\_per\_gallon has an extreme value 1000 that is an outlier in the distribution

and definitely is a misleading data that needs some adjustment.

Summary of Miles per gallon:

```
count      406.000000
mean       42.755665
std        136.102120
min         9.000000
25%        17.500000
50%        23.000000
75%        29.800000
max        1000.000000
Name: Miles per gallon, dtype: float64
```

2. The second attribute that requires attention is the "Horsepower". Since horsepower is a ratio,

0 horsepower means absence of horsepower that does not make sense in the real world. Therefore,

0's need some other relevant values.


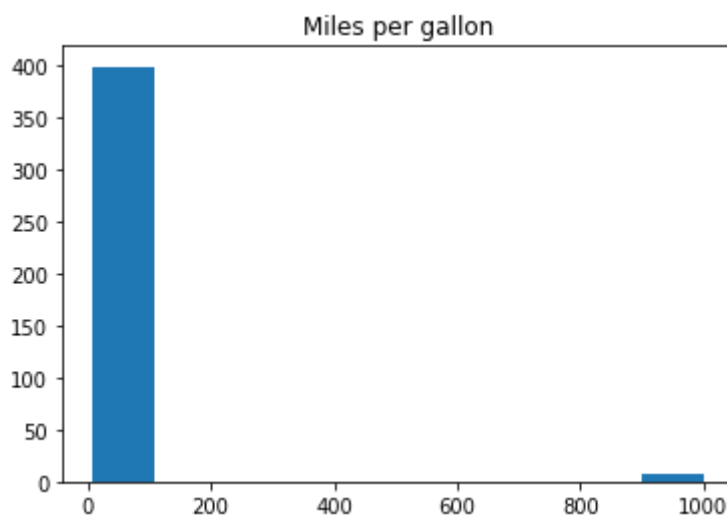
```
count      405.000000
mean       103.785185
std        40.241668
min         0.000000
25%        75.000000
50%        94.000000
75%       129.000000
max       230.000000
Name: Horsepower, dtype: float64
```




Other than those two attributes, the rest of the attributes does not have misleading datapoints.




"""

```
/Library/Python/3.7/site-packages/numpy/lib/histograms.py:839: RuntimeWarning: invalid value encountered in greater_equal
    keep = (tmp_a >= first_edge)
/Library/Python/3.7/site-packages/numpy/lib/histograms.py:840: RuntimeWarning: invalid value encountered in less_equal
    keep &= (tmp_a <= last_edge)
```

```
Out[2]: ' Answer to question number 2\n Here from seeing the distribution in the histograms, some misleading datapoints are observed.\n    The misleading data points are discussed below:\n    1. The first attribute Miles_per_gallon has an extreme value 1000 that is an outlier in the distribution\n    and definately is a misleading data that needs some adjustment. Summary of Miles per gallon:\n\n    \ncount      406.000000\nmean      42.755665\nstd      136.102120\nmin       9.000000\n25%      17.500000\n50%      23.000000\n75%      29.800000\nmax      1000.000000\nName: Miles per gallon, dtype: float64\n\n\n\n    2. The second attribute that requires attention is the "Horsepower". Since horsepower is a ratio,\n    0 horsepower means absence of horsepower that does not make sense in the real world. Therefore, \n    0\'s need some other relevant values.\n\n    \ncount      405.000000\nmean      103.785185\nstd      40.241668\nmin       0.000000\n25%      75.000000\n50%      94.000000\n75%      129.000000\nmax      230.000000\nName: Horsepower, dtype: float64\n\n\n    Other than those two attributes, the rest of the attributes does not have misleading datapoints. \n    '
```







```
In [3]: # Lets change all the values equal to 1000 for miles_Per_gallon to NaN
auto_Mpg["Miles per gallon"] = auto_Mpg["Miles per gallon"].replace({1000:np.NaN})
print(f" The new distribution of Miles Per Gallon is: ")
auto_Mpg["Miles per gallon"].describe()
```

The new distribution of Miles Per Gallon is:

```
Out[3]: count      398.000000
mean        23.514573
std         7.815984
min         9.000000
25%        17.500000
50%        23.000000
75%        29.000000
max        46.600000
Name: Miles per gallon, dtype: float64
```

```
In [4]: # Again, lets change all the values equal to 0 for horsepower to NaN
auto_Mpg["Horsepower"] = auto_Mpg["Horsepower"].replace({0:np.NaN})
print(f" The new distribution of Horsepower is: ")
auto_Mpg["Horsepower"].describe()
```

The new distribution of Horsepower is:

```
Out[4]: count      400.000000
mean       105.082500
std        38.768779
min        46.000000
25%        75.750000
50%        95.000000
75%       130.000000
max       230.000000
Name: Horsepower, dtype: float64
```

```
In [5]: # Now lets create a table with only three attributes (miles_per_gallon,
         Cylinders, Horsepower)
new_Table = auto_Mpg[["Miles per gallon", "Cylinders", "Horsepower"]]
new_Table.describe()
```

```
Out[5]:
```

	Miles per gallon	Cylinders	Horsepower
<b>count</b>	398.000000	406.000000	400.000000
<b>mean</b>	23.514573	5.475369	105.082500
<b>std</b>	7.815984	1.712160	38.768779
<b>min</b>	9.000000	3.000000	46.000000
<b>25%</b>	17.500000	4.000000	75.750000
<b>50%</b>	23.000000	4.000000	95.000000
<b>75%</b>	29.000000	8.000000	130.000000
<b>max</b>	46.600000	8.000000	230.000000

In [ ]:

```
In [6]: # Now lets find an appropriate value to replace the NaN's in the miles_p
er_gallon column
        """ One approach is to find cylinders for each cars is to find a mean to
        replace those NaNs.
           We are going to take this approach to fill in the blanks or NaNs."""

        # All rows with NaN mileage

new_Table[new_Table['Miles per gallon'].isnull()]
```

Out[6]:

	Miles per gallon	Cylinders	Horsepower
10	NaN	4.0	115.0
11	NaN	8.0	165.0
12	NaN	8.0	153.0
13	NaN	8.0	175.0
14	NaN	8.0	175.0
17	NaN	8.0	140.0
39	NaN	4.0	48.0
367	NaN	4.0	110.0

```
In [7]: # All rows with NaN horsepowers
new_Table[new_Table['Horsepower'].isnull()]
```

Out[7]:

	Miles per gallon	Cylinders	Horsepower
38	25.0	4.0	NaN
133	21.0	6.0	NaN
337	40.9	4.0	NaN
343	23.6	4.0	NaN
361	34.5	4.0	NaN
382	23.0	4.0	NaN



```
In [8]: # Calculating the mean mileage for different cars
def meanOfMileage(numOfCylinders):
    num_Cylinder_Cars = new_Table[new_Table.Cylinders == numOfCylinders]
    meanOfMileage = (num_Cylinder_Cars['Miles per gallon']).mean()
    return meanOfMileage

# Function to calculate the mean of horsepower for different cars
def meanOfHorsepower(numOfCylinders):
    num_Cylinder_Cars = new_Table[new_Table.Cylinders == numOfCylinders]
    meanOfHorsepower = (num_Cylinder_Cars['Horsepower']).mean()
    return meanOfHorsepower

# Function to update the horsepower
def horsepowerUpdater(numOfCylinders):
    num_Cylinder_Cars = new_Table[new_Table.Cylinders == numOfCylinders]
    NaNHorsepower = num_Cylinder_Cars[num_Cylinder_Cars['Horsepower'].is
null()]
    NaNHorsepower['Horsepower'] = meanOfHorsepower(numOfCylinders)
    return NaNHorsepower

#function to update the mileage
def mileageUpdater(numOfCylinders):
    num_Cylinder_Cars = new_Table[new_Table.Cylinders == numOfCylinders]
    NaNMileage = num_Cylinder_Cars[num_Cylinder_Cars['Miles per gallon']
.isnull()]
    NaNMileage['Miles per gallon'] = meanOfMileage(numOfCylinders)
    return NaNMileage

updatedHorsepower4 = horsepowerUpdater(4)
updatedHorsepower6 = horsepowerUpdater(6)
updatedHorsepower8 = horsepowerUpdater(8)
updatedHorsepower = pd.concat([updatedHorsepower4, updatedHorsepower6, upd
atedHorsepower8])
updatedHorsepower
```

/Library/Python/3.7/site-packages/ipykernel\_launcher.py:19: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

Out[8]:

	Miles per gallon	Cylinders	Horsepower
<b>38</b>	25.0	4.0	78.470297
<b>337</b>	40.9	4.0	78.470297
<b>343</b>	23.6	4.0	78.470297
<b>361</b>	34.5	4.0	78.470297
<b>382</b>	23.0	4.0	78.470297
<b>133</b>	21.0	6.0	101.506024

```
In [9]: updatedMileage4 = mileageUpdater(4)
updatedMileage6 = mileageUpdater(6)
updatedMileage8 = mileageUpdater(8)
updatedMileage = pd.concat([updatedMileage4,updatedMileage6,updatedMileage8])
updatedMileage
```

/Library/Python/3.7/site-packages/ipykernel\_launcher.py:27: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

Out[9]:

	Miles per gallon	Cylinders	Horsepower
<b>10</b>	29.286765	4.0	115.0
<b>39</b>	29.286765	4.0	48.0
<b>367</b>	29.286765	4.0	110.0
<b>11</b>	14.963107	8.0	165.0
<b>12</b>	14.963107	8.0	153.0
<b>13</b>	14.963107	8.0	175.0
<b>14</b>	14.963107	8.0	175.0
<b>17</b>	14.963107	8.0	140.0

```
In [10]: # Add two dataframes together
updatedVals = pd.concat([updatedMileage,updatedHorsepower])
print("The dataframe with updated values are: ")
updatedVals
```

The dataframe with updated values are:

Out[10]:

	Miles per gallon	Cylinders	Horsepower
<b>10</b>	29.286765	4.0	115.000000
<b>39</b>	29.286765	4.0	48.000000
<b>367</b>	29.286765	4.0	110.000000
<b>11</b>	14.963107	8.0	165.000000
<b>12</b>	14.963107	8.0	153.000000
<b>13</b>	14.963107	8.0	175.000000
<b>14</b>	14.963107	8.0	175.000000
<b>17</b>	14.963107	8.0	140.000000
<b>38</b>	25.000000	4.0	78.470297
<b>337</b>	40.900000	4.0	78.470297
<b>343</b>	23.600000	4.0	78.470297
<b>361</b>	34.500000	4.0	78.470297
<b>382</b>	23.000000	4.0	78.470297
<b>133</b>	21.000000	6.0	101.506024

```
In [11]: # Updating the original dataset with the interpolated values
auto_Mpg.update(updatedVals, join='left', overwrite=True, filter_func=None, errors='ignore')
print("The new dataset is: ")
auto_Mpg.head()
```

The new dataset is:

Out[11]:

	Miles per gallon	Cylinders	Displacement	Horsepower	Weight	Acceleration	Model year	Origin	Car name
0	18.0	8.0	307.0	130.0	3504.0	12.0	70.0	1.0	chevrolet chevelle malibu
1	15.0	8.0	350.0	165.0	3693.0	11.5	70.0	1.0	buick skylark 320
2	18.0	8.0	318.0	150.0	3436.0	11.0	70.0	1.0	plymouth satellite
3	16.0	8.0	304.0	150.0	3433.0	12.0	70.0	1.0	amc rebel sst
4	17.0	8.0	302.0	140.0	3449.0	10.5	70.0	1.0	ford torino

```
In [12]: # Plotting histograms with cleaned dataset
# Plotting all the histograms
# histogram for miles per gallon
plt.figure()
plt.hist(auto_Mpg["Miles per gallon"])
plt.title("Miles per gallon")

#histogram for Cylinders
plt.figure()
plt.hist(auto_Mpg["Cylinders"])
plt.title("Cylinders")

#histogram for Displacement
plt.figure()
plt.hist(auto_Mpg["Displacement"])
plt.title("Displacement")

# histogram for Horsepower
plt.figure()
plt.hist(auto_Mpg["Horsepower"])
plt.title("Horsepower")

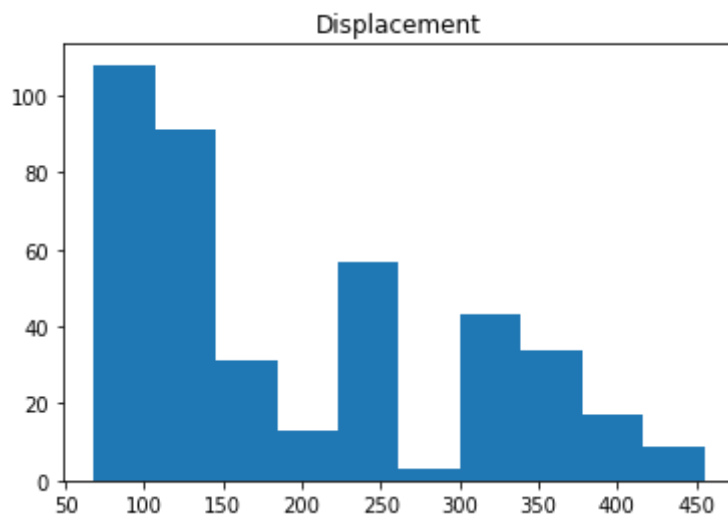


#histogram for Weight
plt.figure()
plt.hist(auto_Mpg["Weight"])
plt.title("Weight")




#histogram for Acceleration
plt.figure()
plt.hist(auto_Mpg["Acceleration"])
plt.title("Acceleration")


#histogram for Model Year
plt.figure()
plt.hist(auto_Mpg["Model year"])
plt.title("Model year")

#histogram for origin
plt.figure()
plt.hist(auto_Mpg["Origin"])
plt.title("Origin")
```

```
Out[12]: Text(0.5, 1.0, 'Origin')
```







```
In [13]: # Exporting the file as a .csv to local folder
auto_Mpg.to_csv("Auto_mpg_adjust.csv", index = False)
```

## Part 3

```
In [14]: # Problem 3
        """ Using z values to check for outliers in the dataset """
        from scipy import stats
        auto_Mpg['milesPerGallon_Z'] = stats.zscore(auto_Mpg["Miles per gallon"])
        # print(auto_Mpg['Miles per gallon_Z'])

        # Looking for outliers in the dataset using z values
        MPGOutliers = auto_Mpg.query('milesPerGallon_Z >3 | milesPerGallon_Z <-3')
        numberOfMPGOutliers = len(MPGOutliers)
        print(f'There are %d outliers present in MPG data.' % numberOfMPGOutliers)
```

There are 0 outliers present in MPG data.



```
In [15]: # Looking into Cylinder data
auto_Mpg['cylinders_Z'] = stats.zscore(auto_Mpg["Cylinders"])
# print(auto_Mpg['cylinders_Z'])
cylindersOutliers = auto_Mpg.query('cylinders_Z >3 | cylinders_Z <-3')
numberOfCylindersOutliers = len(cylindersOutliers)
print(f'There are %d outliers present in Cylinders data.' % numberOfCylindersOutliers)
```

There are 0 outliers present in Cylinders data.

```
In [16]: # Looking into Displacement data
auto_Mpg['displacement_Z'] = stats.zscore(auto_Mpg["Displacement"])
# print(auto_Mpg['displacement_Z'])
displacementOutliers = auto_Mpg.query('displacement_Z >3 | displacement_Z <-3')
numberOfDisplacementOutliers = len(displacementOutliers)
print(f'There are %d outliers present in Displacements data.' % numberOfDisplacementOutliers)
```

There are 0 outliers present in Displacements data.

```
In [17]: # Looking into Horsepower data
auto_Mpg['horsepower_Z'] = stats.zscore(auto_Mpg["Horsepower"])
# print(auto_Mpg['horsepower_Z'])
horsepowerOutliers = auto_Mpg.query('horsepower_Z >3 | horsepower_Z <-3')
numberOfHorsepowerOutliers = len(horsepowerOutliers)
print(f'There are %d outliers present in Horsepower data.' % numberOfHorsepowerOutliers)
horsepowerSort = auto_Mpg.sort_values('horsepower_Z', ascending = False)
horsepowerSort.head(n = 4)
```

There are 4 outliers present in Horsepower data.

Out[17]:

	Miles per gallon	Cylinders	Displacement	Horsepower	Weight	Acceleration	Model year	Origin	Car name
123	16.0	8.0	400.0	230.0	4278.0	9.5	73.0	1.0	pontiac grand prix
102	12.0	8.0	455.0	225.0	4951.0	11.0	73.0	1.0	buick electra 225 custom
19	14.0	8.0	455.0	225.0	3086.0	10.0	70.0	1.0	buick estate wagon (sw)
8	14.0	8.0	455.0	225.0	4425.0	10.0	70.0	1.0	pontiac catalina

```
In [18]: # Looking into Weight data
auto_Mpg['weight_Z'] = stats.zscore(auto_Mpg["Weight"])
# print(auto_Mpg['weight_Z'])
weightOutliers = auto_Mpg.query('weight_Z >3 | weight_Z <-3')
numberOfWeightOutliers = len(weightOutliers)
print(f'There are %d outliers present in Weight data.' % numberOfWeightOutliers)
```

There are 0 outliers present in Weight data.

```
In [19]: # Looking into Acceleration data
auto_Mpg['acceleration_Z'] = stats.zscore(auto_Mpg["Acceleration"])
# print(auto_Mpg['acceleration_Z'])
accelerationOutliers = auto_Mpg.query('acceleration_Z >3 | acceleration_Z <-3')
numberOfAccelerationOutliers = len(accelerationOutliers)
print(f'There are %d outliers present in Acceleration data.' % numberOfAccelerationOutliers)
accelerationSort = auto_Mpg.sort_values('acceleration_Z', ascending = False)
accelerationSort.head(n = 2)
```

There are 2 outliers present in Acceleration data.

Out[19]:

	Miles per gallon	Cylinders	Displacement	Horsepower	Weight	Acceleration	Model year	Origin	Car name
306	27.2	4.0	141.0	71.0	3190.0	24.8	79.0	2.0	peugeot 504
402	44.0	4.0	97.0	52.0	2130.0	24.6	82.0	2.0	vw pickup

```
In [20]: # Looking into modelYear data
auto_Mpg['modelYear_Z'] = stats.zscore(auto_Mpg["Model year"])
# print(auto_Mpg['modelYear_Z'])
modelYearOutliers = auto_Mpg.query('modelYear_Z >3 | modelYear_Z <-3')
numberOfModelYearOutliers = len(modelYearOutliers)
print(f'There are %d outliers present in Model year data.' % numberOfModelYearOutliers)
```

There are 0 outliers present in Model year data.

```
In [21]: # Looking into Origin data
auto_Mpg['origin_Z'] = stats.zscore(auto_Mpg["Origin"])
# print(auto_Mpg['origin_Z'])
originOutliers = auto_Mpg.query('origin_Z >3 | origin_Z <-3')
numberOfOriginOutliers = len(originOutliers)
print(f'There are %d outliers present in Origin data.' % numberOfOriginOutliers)
```

There are 0 outliers present in Origin data.

## Part 4

```
In [22]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
pd.set_option('display.max_rows', 500)

cereals = pd.read_csv("cereals.csv")
cereals.head()
```

Out[22]:

	Name	Manuf	Type	Calories	Protein	Fat	Sodium	Fiber	Carbo	Sugars	...
0	100%_Bran	N	C	70	4	1	130	10.0	5.0	6.0	...
1	100%_Natural_Bran	Q	C	120	3	5	15	2.0	8.0	8.0	...
2	All-Bran	K	C	70	4	1	260	9.0	7.0	5.0	...
3	All-Bran_with_Extra_Fiber	K	C	50	4	0	140	14.0	8.0	0.0	...
4	Almond_Delight	R	C	110	2	2	200	1.0	14.0	8.0	...

5 rows × 23 columns


## Part a

```
In [23]: # create a bar graph and normalized bar graph of the "Manuf" variable with "Type" overlay
crosstab_01 = pd.crosstab(cereals["Manuf"],cereals['Type'])
plt.figure()
crosstab_01.plot(kind='bar', stacked=True)
plt.title('Bar Graph of Manuf with Type Overlay')

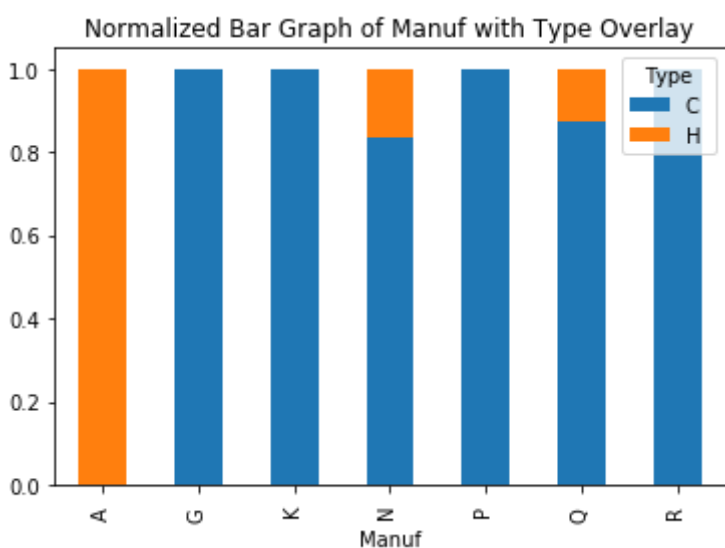
# Normalized bar graph
plt.figure()
crosstab_norm = crosstab_01.div(crosstab_01.sum(1),axis=0)
crosstab_norm.plot(kind='bar', stacked=True)
plt.title('Normalized Bar Graph of Manuf with Type Overlay')
```

Out[23]: Text(0.5, 1.0, 'Normalized Bar Graph of Manuf with Type Overlay')

<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>



## Part b

```
In [24]: # create a contingency table of "Manuf" and "Type"

crosstab_02 = pd.crosstab(cereals["Type"],cereals['Manuf'])
crosstab_norm_02 = round(crosstab_02.div(crosstab_02.sum(0),axis=1)* 100
,1)
crosstab_02
```

Out[24]:

Manuf	A	G	K	N	P	Q	R
Type							
C	0	22	23	5	9	7	8
H	1	0	0	1	0	1	0

```
In [25]: crosstab_norm_02
```

Out[25]:

Manuf	A	G	K	N	P	Q	R
Type							
C	0.0	100.0	100.0	83.3	100.0	87.5	100.0
H	100.0	0.0	0.0	16.7	0.0	12.5	0.0

## Part c

```

In [26]: # create a histogram and normalized histogram of "Calories" with "Manuf"
         overlay
crosstab_03 = pd.crosstab(cereals["Calories"],cereals['Manuf'])
plt.figure()
crosstab_03.plot(kind='bar', stacked=True)
plt.title('Bar Graph of Calories with Manuf Overlay')

# Normalized bar graph
plt.figure()
crosstab_03_norm = crosstab_03.div(crosstab_03.sum(1),axis=0)
crosstab_03_norm.plot(kind='bar', stacked=True)
plt.title('Normalized Bar Graph of Calories with Manuf Overlay')

# Stacked histogram
plt.figure()
cerealCaloriesManuf_A = cereals[cereals.Manuf == 'A ']['Calories']
cerealCaloriesManuf_G = cereals[cereals.Manuf == "G "]['Calories']
cerealCaloriesManuf_K = cereals[cereals.Manuf == "K "]['Calories']
cerealCaloriesManuf_N = cereals[cereals.Manuf == "N "]['Calories']
cerealCaloriesManuf_P = cereals[cereals.Manuf == "P "]['Calories']
cerealCaloriesManuf_Q = cereals[cereals.Manuf == "Q "]['Calories']
cerealCaloriesManuf_R = cereals[cereals.Manuf == "R "]['Calories']
plt.hist([cerealCaloriesManuf_A,cerealCaloriesManuf_G,cerealCaloriesManu
f_K,
          cerealCaloriesManuf_N,cerealCaloriesManuf_P,cerealCaloriesManu
f_Q,
          cerealCaloriesManuf_R],bins = 10, stacked = True)
plt.legend(['Manuf A', 'Manuf G', 'Manuf K', 'Manuf N', 'Manuf P', 'Manu
f Q', 'Manuf R'])
plt.title('Histogram of Calories with Manuf Overlay')
plt.ylabel('Frequency')
plt.xlabel('Calories')

# Normalized histogram of Calories with Manuf Overlay
(n,bins,patches)= plt.hist([cerealCaloriesManuf_A,cerealCaloriesManuf_G,
cerealCaloriesManuf_K,
          cerealCaloriesManuf_N,cerealCaloriesManuf_P,cerealCaloriesManu
f_Q,
          cerealCaloriesManuf_R],bins = 10, stacked = True)

# Creating a new plot
plt.figure()
n_table = np.column_stack((n[0],n[1]))
n_norm = n_table/n_table.sum(axis = 1)[: ,None]

# creating an array of bin cuts
our_bins = np.column_stack((bins[0:10],bins[1:11]))
p1 = plt.bar(x = our_bins[:,0],height = n_norm[:,0],width = our_bins[:,1
]-our_bins[:,0])
p2 = plt.bar(x = our_bins[:,0],height = n_norm[:,1],width = our_bins[:,1
]-our_bins[:,0],bottom = n_norm[:,0])


```

```
''' I could not figure this out. '''
```


```
/Library/Python/3.7/site-packages/ipykernel_launcher.py:41: RuntimeWarning: invalid value encountered in true_divide
```

```
Out[26]: ' I could not figure this out. '
```


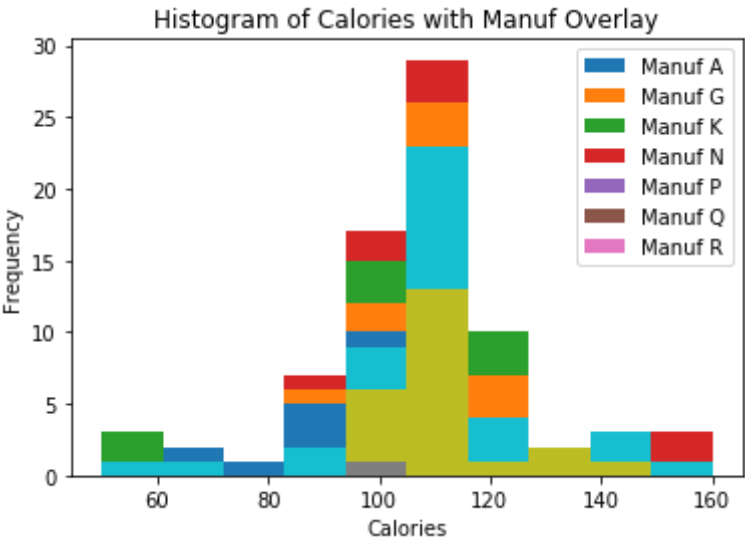
```
<Figure size 432x288 with 0 Axes>
```



```
<Figure size 432x288 with 0 Axes>
```







Part d

```
In [27]: # Bin the "Calories" variable using bins for 0 to 90, 91-110, and over 1
10 calories and ...
# create a bar chart of the binned calories variable with "Manuf" overla
y
# Binning the calories
cereals["caloriesBinned"] = pd.cut(x = cereals['Calories'],bins = [0,91,
111,165],
                                labels=['0-90','90-110','Over 110'],
right = False)
crosstab_04 = pd.crosstab(cereals['caloriesBinned'],cereals['Manuf'])
crosstab_04.plot(kind='bar',stacked = True,title = 'Bar chart of Binned
Caloris with Manuf Overlay')
```

Out[27]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1229259e8>

